

Time Series Forecasting_Predicting Sales using ARIMA

```
In [1]: #https://www.youtube.com/watch?v=MVsKaYzEggY&list=PLmPJQXJiMoUVr07-VnwDiki89Dgyu
#Changing working directory
import os
print(os.getcwd())
import warnings
warnings.filterwarnings('ignore')
```

```
/Users/pragatigupta/Documents/AI And ML/2023/Time Series/Time Series forecasting
sales
```

```
In [3]: from IPython.display import Image
Image(filename='/Users/pragatigupta/Documents/AI And ML/2023/Time Series/Time Se
```

Out[3]:



Problem: Forecasting Sales

Step1:

Gathered (Extract) Data from all the sources

```
In [368... # Data and package Import
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.arima_model import ARMA
TempData = pd.read_csv('/Users/pragatigupta/Documents/AI And ML/2023/Time Series
TempData.head(5)
```

Out [368...

	Date	Symbol	Series	Prev Close	Open	High	Low	Last	Close	VWAP	Volume
0	2000-01-11	HCLTECH	EQ	580.00	1550.0	1725.00	1492.00	1560.00	1554.45	1582.72	1192200
1	2000-01-12	HCLTECH	EQ	1554.45	1560.0	1678.85	1560.00	1678.85	1678.85	1657.05	344850
2	2000-01-13	HCLTECH	EQ	1678.85	1790.0	1813.20	1781.00	1813.20	1813.20	1804.69	53000
3	2000-01-14	HCLTECH	EQ	1813.20	1958.3	1958.30	1835.00	1958.30	1958.30	1939.90	270950
4	2000-01-17	HCLTECH	EQ	1958.30	2115.0	2115.00	1801.65	1801.65	1801.65	1990.55	428800

Step2:Data Preprocessing:

Clean and preprocess the historical sales data, handling missing values and outliers if necessary. Explore the data visually to identify any apparent trends, seasonality, or irregular patterns.

In [369...

```
#Data Cleaning
Sales_Data = TempData.dropna()
```

In [370...

```
Sales_Data=Sales_Data.drop_duplicates(subset=['Date'])
```

In [371...

```
Sales_Data.index = pd.to_datetime(Sales_Data.Date)
```

In [372...

```
print(Sales_Data.index.min())
print(Sales_Data.index.max())

2011-06-01 00:00:00
2020-11-27 00:00:00
```

In [373...

```
Sales_Data = Sales_Data["Prev Close"][ '2013-01-01': '2013-12-2' ] # select the col
#or other method is to drop unwanted columns
#Sales_Data = Sales_Data.drop(['col1','col2'], axis=1)
Sales_Data.describe()
```

Out [373...

```
count      230.000000
mean       852.953478
std        156.484472
min        618.700000
25%        736.350000
50%        777.450000
75%       1023.962500
max       1161.150000
Name: Prev Close, dtype: float64
```

In [374...

```
Sales_Data.head(10)
```

Out[374...

```
Date
2013-01-01    618.70
2013-01-02    622.15
2013-01-03    625.25
2013-01-04    625.95
2013-01-07    634.05
2013-01-08    627.90
2013-01-09    634.70
2013-01-10    635.85
2013-01-11    641.60
2013-01-14    644.85
Name: Prev Close, dtype: float64
```


Step 3: Data Exploration

In [375...

```
#Data Exploration
plt.figure(figsize=(15,7))
fig = plt.figure(1)
ax1 = fig.add_subplot(111)
ax1.set_xlabel('Time Frame')
ax1.set_ylabel('Sale Prices')
ax1.plot(Sales_Data)
```

Out[375...

```
[<matplotlib.lines.Line2D at 0x7ff51aea45b0>]
```



Step 4: Feature Engineering:

Create additional features based on external factors like marketing spend, seasonal indicators, and economic variables. Lag features: Include lagged sales data to capture autocorrelation.

In [376...


```
# Checking stationarity
```

```
In [377...
# Method 1 - Rolling Statistics
# Method 2 - Duckey fuller
```

```
In [378...
#Determining rolling statistics
rollmean = Sales_Data.rolling(12).mean()
rollstd = Sales_Data.rolling(12).std()

plt.figure(figsize=(16,7))
fig = plt.figure(1)

#Plot rolling statistics:
orig = plt.plot(Sales_Data, color='blue',label='Original')
mean = plt.plot(rollmean, color='red', label='Rolling Mean')
std = plt.plot(rollstd, color='black', label = 'Rolling Std')
plt.legend(loc='best')
plt.title('Rolling Mean & Standard Deviation')
plt.show(block=False)
```



Making Series Stationary

```
In [379...
#Lets try transformation
plt.figure(figsize=(16,7))
fig = plt.figure(1)
#log , square root transformation ect to make the series stationary
import numpy as np
ts_log = np.log(Sales_Data)
plt.plot(ts_log)
```

```
Out[379... [<matplotlib.lines.Line2D at 0x7ff53504ad60>]
```



In [380...

```
#Decomposition
from statsmodels.tsa.seasonal import seasonal_decompose
decomposition = seasonal_decompose(ts_log,freq=1,model = 'multiplicative')


trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid

plt.figure(figsize=(16,7))
fig = plt.figure(1)

plt.subplot(411)
plt.plot(ts_log, label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend, label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal,label='Seasonality')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual, label='Residuals')
plt.legend(loc='best')
```

Out[380...

```
<matplotlib.legend.Legend at 0x7ff57252d0a0>
```




In [381]...

```
#Lets try differencing
plt.figure(figsize=(16,7))
fig = plt.figure(1)
ts_log_diff = ts_log - ts_log.shift()
plt.plot(ts_log_diff)

#Determining rolling statistics
rollmean = ts_log_diff.rolling(12).mean()
rollstd = ts_log_diff.rolling(12).std()


#Plot rolling statistics:
orig = plt.plot(ts_log_diff, color='blue', label='Original')
mean = plt.plot(rollmean, color='red', label='Rolling Mean')
std = plt.plot(rollstd, color='black', label='Rolling Std')
plt.legend(loc='best')
plt.title('Rolling Mean & Standard Deviation')
plt.show(block=False)
```



```
In [382... Sales_Data.sort_index(inplace= True)
```

```
In [389... from statsmodels.tsa.stattools import acf, pacf
lag_acf = acf(ts_log_diff, nlags=20)
lag_pacf = pacf(ts_log_diff, nlags=20)
```

```
In [392... import statsmodels.api as sm
fig = plt.figure(figsize=(12,8))
ax1 = fig.add_subplot(211)
fig = sm.graphics.tsa.plot_acf(ts_log_diff.dropna(),lags=40,ax=ax1)
ax2 = fig.add_subplot(212)
fig = sm.graphics.tsa.plot_pacf(ts_log_diff.dropna(),lags=40,ax=ax2)
#Chart: helps to get correct AR and MA values
# highlighted part = confidence interval
# first line which touch or cross the highlighted part is the correct value
# from the fig AR=2, MA=2
```



```
In [393... from statsmodels.tsa.arima_model import ARIMA
```

```
In [394... type(ts_log_diff)
```

```
Out[394... pandas.core.series.Series
```

```
In [395... #ts_log_diff.dropna()
ts_log_diff = ts_log_diff[~ts_log_diff.isnull()]
```

In [396...

```
plt.figure(figsize=(16,8))
#ts_log_diff.dropna(inplace=True)
model = ARIMA(ts_log_diff, order=(2,1,2))
results_ARIMA = model.fit()
plt.plot(ts_log_diff)
plt.plot(results_ARIMA.fittedvalues, color='red')
```

/Users/pragatigupta/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no'

/Users/pragatigupta/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no'

This problem is unconstrained.

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 5 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= -2.55431D+00 |proj g|= 2.26589D+00

At iterate 5 f= -2.55431D+00 |proj g|= 4.08178D-01

At iterate 10 f= -2.55432D+00 |proj g|= 6.35795D+00

At iterate 15 f= -2.55449D+00 |proj g|= 5.80723D+01

At iterate 20 f= -2.55492D+00 |proj g|= 3.62180D-02

At iterate 25 f= -2.55492D+00 |proj g|= 5.25044D-01

At iterate 30 f= -2.55493D+00 |proj g|= 7.10730D+00

At iterate 35 f= -2.55514D+00 |proj g|= 2.78026D+01

At iterate 40 f= -2.55518D+00 |proj g|= 2.52323D-03

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

Skip = number of BFGS updates skipped

Nact = number of active bounds at final generalized Cauchy point

Projg = norm of the final projected gradient

F = final function value

* * *

N	Tit	Tnf	Tnint	Skip	Nact	Projg	F
5	41	59	1	0	0	2.523D-03	-2.555D+00


F = -2.5551832218989681

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

/Users/pragatigupta/opt/anaconda3/lib/python3.9/site-packages/statsmodels/base/odel.py:547: HessianInversionWarning: Inverting hessian failed, no bse or cov_params available

warnings.warn('Inverting hessian failed, no bse or cov_params '

Out[396... [



Taking results back to original scale

In [397... `ARIMA_diff_predictions = pd.Series(results_ARIMA.fittedvalues, copy=True)`
`print(ARIMA_diff_predictions.head())`

```
Date
2013-01-03    -0.000005
2013-01-04     0.000273
2013-01-07     0.002712
2013-01-08    -0.006513
2013-01-09     0.011963
dtype: float64
```

In [398... `ARIMA_diff_predictions_cumsum = ARIMA_diff_predictions.cumsum()`
`print(ARIMA_diff_predictions_cumsum.head())`
reverse the log : cumulative sum to difference the time series

```
Date
2013-01-03    -0.000005
2013-01-04     0.000268
2013-01-07     0.002980
2013-01-08    -0.003533
2013-01-09     0.008430
dtype: float64
```

In [399... `ARIMA_log_prediction = pd.Series(ts_log.iloc[0], index=ts_log.index)`
`ARIMA_log_prediction = ARIMA_log_prediction.add(ARIMA_diff_predictions_cumsum, fi`

```
ARIMA_log_prediction.head()
# reverse the diffrese
```

```
Out[399... Date
2013-01-01    6.427621
2013-01-02    6.427621
2013-01-03    6.427615
2013-01-04    6.427888
2013-01-07    6.430600
dtype: float64
```

```
In [400... plt.figure(figsize=(12,8))
predictions_ARIMA = np.exp(ARIMA_log_prediction)
plt.plot(Sales_Data)
plt.plot(predictions_ARIMA)
plt.title('RMSE: %.4f'% np.sqrt(sum((predictions_ARIMA-Sales_Data)**2)/len(Sales
# reverse the log: expontital
```

```
Out[400... Text(0.5, 1.0, 'RMSE: 266.6637')
```



```
In [401... results_ARIMA.predict(10,20)
```

```
Out[401... Date
2013-01-16    -0.032051
2013-01-17     0.016081
2013-01-18    -0.003840
2013-01-21    -0.037005
2013-01-22     0.008956
2013-01-23    -0.005627
2013-01-24     0.033022
2013-01-25     0.035053
2013-01-28     0.020183
```

```
2013-01-29    -0.017208
2013-01-30    -0.016286
dtype: float64
```

In [402...

```
#Auto ARIMA: pip install pmdarima on command window
```

In [403...

```
import pmdarima as pm
def arimamodel(timeseries):
    automodel = pm.auto_arima(timeseries,
                              start_p=3,
                              start_q=3,
                              max_p=5,
                              max_q=5,
                              test="adf",
                              seasonal=True,
                              trace=True)

    return automodel
```

In [404...

```
arimamodel(ts_log)
```

Performing stepwise search to minimize aic

```
ARIMA(3,1,3)(0,0,0)[0] intercept    : AIC=-1159.043, Time=0.31 sec
ARIMA(0,1,0)(0,0,0)[0] intercept    : AIC=-1167.424, Time=0.02 sec
ARIMA(1,1,0)(0,0,0)[0] intercept    : AIC=-1165.456, Time=0.03 sec
ARIMA(0,1,1)(0,0,0)[0] intercept    : AIC=-1165.465, Time=0.04 sec
ARIMA(0,1,0)(0,0,0)[0]              : AIC=-1165.525, Time=0.02 sec
ARIMA(1,1,1)(0,0,0)[0] intercept    : AIC=-1163.485, Time=0.08 sec
```

Best model: ARIMA(0,1,0)(0,0,0)[0] intercept

Total fit time: 0.518 seconds

Out[404...

```
ARIMA(order=(0, 1, 0), scoring_args={}, suppress_warnings=True)
```

In [411...

```
plt.figure(figsize=(16,8))
#ts_log_diff.dropna(inplace=True)
model = ARIMA(ts_log_diff, order=(0,1,0))
results_ARIMA = model.fit()
plt.plot(ts_log_diff)
plt.plot(results_ARIMA.fittedvalues, color='red')
```

RUNNING THE L-BFGS-B CODE

* * *

Machine precision = 2.220D-16

N = 1 M = 12

At X0 0 variables are exactly at the bounds

At iterate 0 f= -2.21608D+00 |proj g|= 7.19425D-06

* * *

Tit = total number of iterations

Tnf = total number of function evaluations

Tnint = total number of segments explored during Cauchy searches

```

Skip = number of BFGS updates skipped
Nact = number of active bounds at final generalized Cauchy point
Projg = norm of the final projected gradient
F      = final function value

```

* * *

```

      N      Tit      Tnf  Tnint  Skip  Nact      Projg      F
      1       1      14     1     0     0    6.839D-06  -2.216D+00
F = -2.2160806223129810

```

CONVERGENCE: REL_REDUCTION_OF_F_<=_FACTR*EPSMCH

/Users/pragatigupta/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no'


/Users/pragatigupta/opt/anaconda3/lib/python3.9/site-packages/statsmodels/tsa/base/tsa_model.py:581: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

warnings.warn('A date index has been provided, but it has no'

This problem is unconstrained.

Warning: more than 10 function and gradient evaluations in the last line search. Termination may possibly be caused by a bad search direction.

Out[411... [



In [412...

```

ARIMA_diff_predictions = pd.Series(results_ARIMA.fittedvalues, copy=True)
print(ARIMA_diff_predictions.head())

```

```

Date
2013-01-03    0.000073
2013-01-04    0.000073
2013-01-07    0.000073
2013-01-08    0.000073
2013-01-09    0.000073
dtype: float64

```

In [413...

```
ARIMA_diff_predictions_cumsum = ARIMA_diff_predictions.cumsum()
print(ARIMA_diff_predictions_cumsum.head())
# reverse the log : cumulatevive sum to diffrence the time series
```

```
Date
2013-01-03    0.000073
2013-01-04    0.000146
2013-01-07    0.000219
2013-01-08    0.000292
2013-01-09    0.000365
dtype: float64
```

In [414...

```
ARIMA_log_prediction = pd.Series(ts_log.iloc[0], index=ts_log.index)
ARIMA_log_prediction = ARIMA_log_prediction.add(ARIMA_diff_predictions_cumsum, fi
ARIMA_log_prediction.head()
# reverse the diffrese
```

Out[414...


```
Date
2013-01-01    6.427621
2013-01-02    6.427621
2013-01-03    6.427694
2013-01-04    6.427767
2013-01-07    6.427840
dtype: float64
```

In [415...

```
plt.figure(figsize=(12,8))
predictions_ARIMA = np.exp(ARIMA_log_prediction)
plt.plot(Sales_Data)
plt.plot(predictions_ARIMA)
plt.title('RMSE: %.4f'% np.sqrt(sum((predictions_ARIMA-Sales_Data)**2)/len(Sales
# reverse the log: expontital
```

Out[415...

```
Text(0.5, 1.0, 'RMSE: 275.6730')
```



In []:

In []: